

# DSC 680 Project 1

December 18, 2025

```
[3]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_curve, roc_auc_score
```

```
[4]: #load flat file dataset & select columns needed for analysis
telecom_data = pd.read_csv('/Volumes/Editing/Bellevue Univ/Masters in Data Science/DSC 680 Applied Data Science/Project 1/WA_Fn-UseC_-Telco-Customer-Churn.csv')
```

```
[5]: telecom_data.head()
```

```
[5]:   customerID  gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService \
0    7590-VHVEG  Female           0      Yes        No         1        No
1    5575-GNVDE    Male           0       No        No        34       Yes
2    3668-QPYBK    Male           0       No        No         2       Yes
3    7795-CFOCW    Male           0       No        No        45        No
4    9237-HQITU  Female           0       No        No         2       Yes

          MultipleLines  InternetService  OnlineSecurity ... DeviceProtection \
0  No phone service            DSL           No     ...        No
1                No            DSL           Yes     ...       Yes
2                No            DSL           Yes     ...        No
3  No phone service            DSL           Yes     ...       Yes
4                No  Fiber optic           No     ...        No

  TechSupport  StreamingTV  StreamingMovies  Contract  PaperlessBilling \
0        No           No           No  Month-to-month        Yes
1        No           No           No      One year        No
2        No           No           No  Month-to-month       Yes
3       Yes           No           No      One year        No
4        No           No           No  Month-to-month       Yes
```

```

          PaymentMethod MonthlyCharges TotalCharges Churn
0      Electronic check        29.85      29.85    No
1      Mailed check         56.95    1889.5    No
2      Mailed check         53.85     108.15   Yes
3 Bank transfer (automatic)     42.30    1840.75    No
4      Electronic check       70.70     151.65   Yes

```

[5 rows x 21 columns]

## 0.1 Data Cleaning

```
[6]: # Count of rows/columns
print(telecom_data.shape)
```

(7043, 21)

```
[7]: telecom_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   customerID      7043 non-null   object 
 1   gender          7043 non-null   object 
 2   SeniorCitizen   7043 non-null   int64  
 3   Partner         7043 non-null   object 
 4   Dependents     7043 non-null   object 
 5   tenure          7043 non-null   int64  
 6   PhoneService    7043 non-null   object 
 7   MultipleLines   7043 non-null   object 
 8   InternetService 7043 non-null   object 
 9   OnlineSecurity  7043 non-null   object 
 10  OnlineBackup    7043 non-null   object 
 11  DeviceProtection 7043 non-null   object 
 12  TechSupport    7043 non-null   object 
 13  StreamingTV    7043 non-null   object 
 14  StreamingMovies 7043 non-null   object 
 15  Contract        7043 non-null   object 
 16  PaperlessBilling 7043 non-null   object 
 17  PaymentMethod   7043 non-null   object 
 18  MonthlyCharges 7043 non-null   float64
 19  TotalCharges   7043 non-null   object 
 20  Churn           7043 non-null   object 

dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

```

```
[8]: telecom_data['Churn'].value_counts()
```

```
[8]: Churn
No      5174
Yes     1869
Name: count, dtype: int64
```

```
[9]: telecom_data['gender'].value_counts()
```

```
[9]: gender
Male     3555
Female   3488
Name: count, dtype: int64
```

```
[10]: telecom_data['Churn'].value_counts(normalize=True)
```

```
[10]: Churn
No      0.73463
Yes     0.26537
Name: proportion, dtype: float64
```

```
[11]: # Convert Total Charges to Numeric data type to determine nulls
telecom_data['TotalCharges'] = pd.to_numeric(telecom_data['TotalCharges'], ↴
    errors='coerce')
telecom_data['MonthlyCharges'] = pd.to_numeric(telecom_data['MonthlyCharges'], ↴
    errors='coerce')
```

```
[12]: # Count of missing values
telecom_data.isnull().sum()
```

```
[12]: customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents     0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV    0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
```

```
MonthlyCharges      0  
TotalCharges       11  
Churn              0  
dtype: int64
```

11 Null values found

```
[13]: # Impute Using tenure x MonthlyCharges  
telecom_data.loc[telecom_data['TotalCharges'].isna(), 'TotalCharges'] = (  
    telecom_data.loc[telecom_data['TotalCharges'].isna(), 'tenure'] *  
    telecom_data.loc[telecom_data['TotalCharges'].isna(), 'MonthlyCharges'])
```

```
[14]: # Final Check to verify missing values  
telecom_data['TotalCharges'].isna().sum()
```

[14]: np.int64(0)

```
[15]: telecom_data.describe
```

```
[15]: <bound method NDFrame.describe of
      Partner Dependents tenure \
0    7590-VHVEG Female          0 Yes No 1
1    5575-GNVDE Male           0 No No 34
2    3668-QPYBK  Male          0 No No 2
3    7795-CFOCW Male           0 No No 45
4    9237-HQITU Female         0 No No 2
...
7038 6840-RESVB Male           0 Yes Yes 24
7039 2234-XADUH Female        0 Yes Yes 72
7040 4801-JZAZL Female        0 Yes Yes 11
7041 8361-LTMKD Male          1 Yes No 4
7042 3186-AJIEK Male           0 No No 66

PhoneService MultipleLines InternetService OnlineSecurity ... \
0           No No phone service DSL No ...
1          Yes No DSL Yes ...
2          Yes No DSL Yes ...
3           No No phone service DSL Yes ...
4          Yes No Fiber optic No ...
...
7038        Yes Yes DSL Yes ...
7039        Yes Yes Fiber optic No ...
7040        No No phone service DSL Yes ...
7041        Yes Yes Fiber optic No ...
7042        Yes No Fiber optic Yes ...

DeviceProtection TechSupport StreamingTV StreamingMovies Contr...
0           No No No No Month-to-mo...
```

```

1           Yes      No      No      No      One year
2           No       No      No      No   Month-to-month
3          Yes      Yes      No      No      One year
4          No       No      No      No   Month-to-month
...
7038        Yes      Yes      Yes      Yes      One year
7039        Yes      No       Yes      Yes      One year
7040        No       No       No      No   Month-to-month
7041        No       No       No      No   Month-to-month
7042        Yes      Yes      Yes      Yes      Two year

    PaperlessBilling      PaymentMethod MonthlyCharges TotalCharges \
0            Yes   Electronic check      29.85      29.85
1            No    Mailed check       56.95    1889.50
2            Yes   Mailed check       53.85     108.15
3            No  Bank transfer (automatic)  42.30    1840.75
4            Yes   Electronic check      70.70     151.65
...
7038        ...      Mailed check      84.80    1990.50
7039        Yes   Credit card (automatic) 103.20    7362.90
7040        Yes   Electronic check      29.60     346.45
7041        Yes   Mailed check       74.40     306.60
7042        Yes  Bank transfer (automatic) 105.65    6844.50

    Churn
0          No
1          No
2         Yes
3          No
4         Yes
...
7038        No
7039        No
7040        No
7041        Yes
7042        No

```

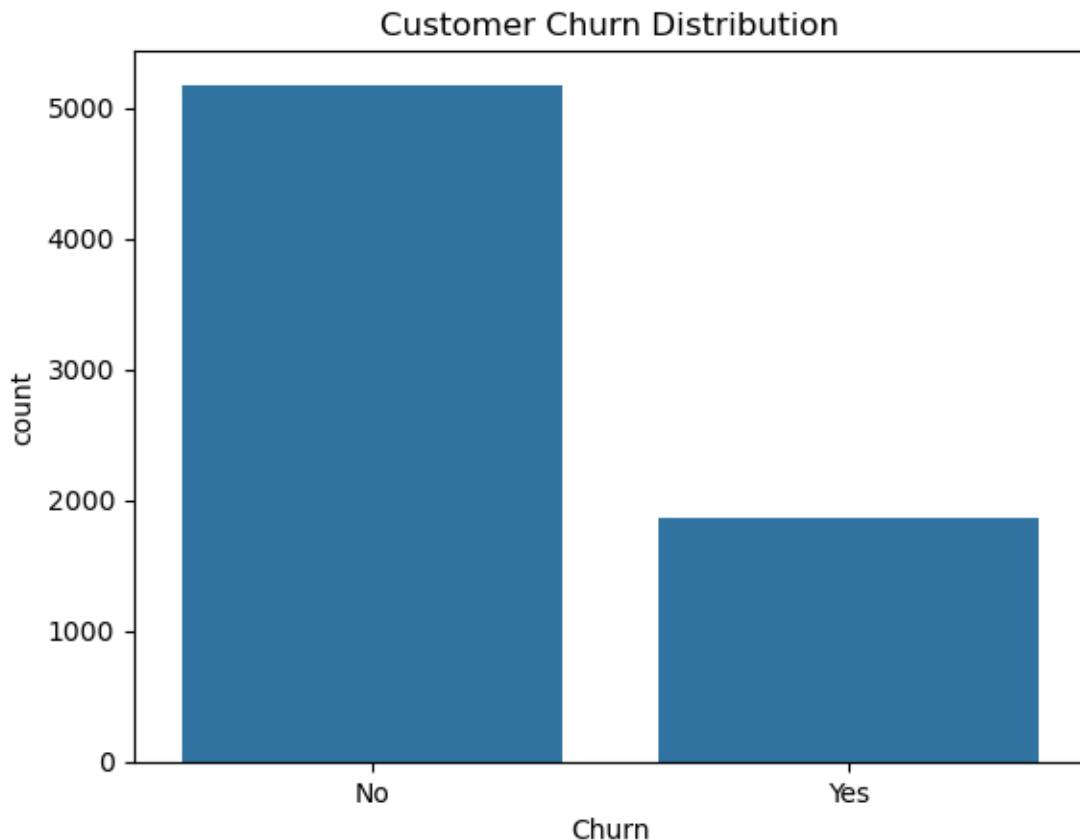
[7043 rows x 21 columns]>

EDA

```
[16]: telecom_data['TotalCharges'] = pd.to_numeric(telecom_data['TotalCharges'],  
    errors='coerce')
```

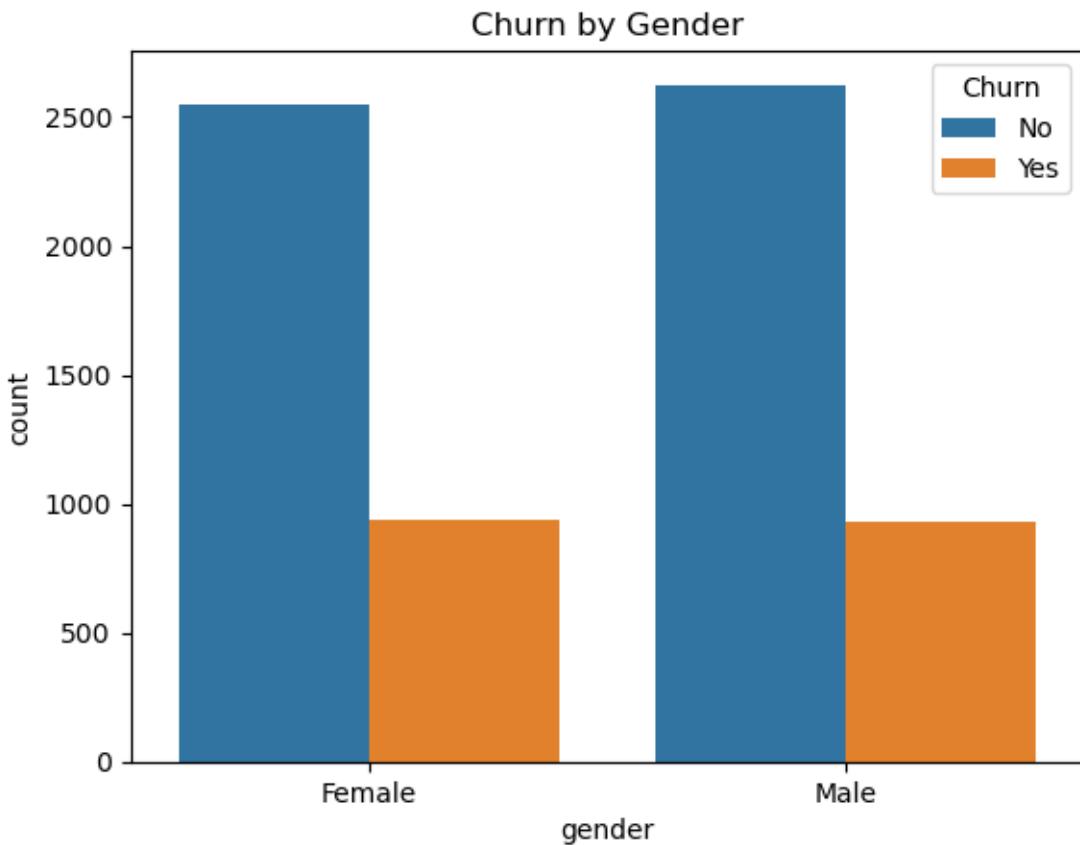
```
[17]: # Drop the customer ID column  
telecom_data.drop('customerID', axis=1, inplace=True)
```

```
[18]: # Visualize churn
sns.countplot(data=telecom_data, x='Churn')
plt.title("Customer Churn Distribution")
plt.show()
```



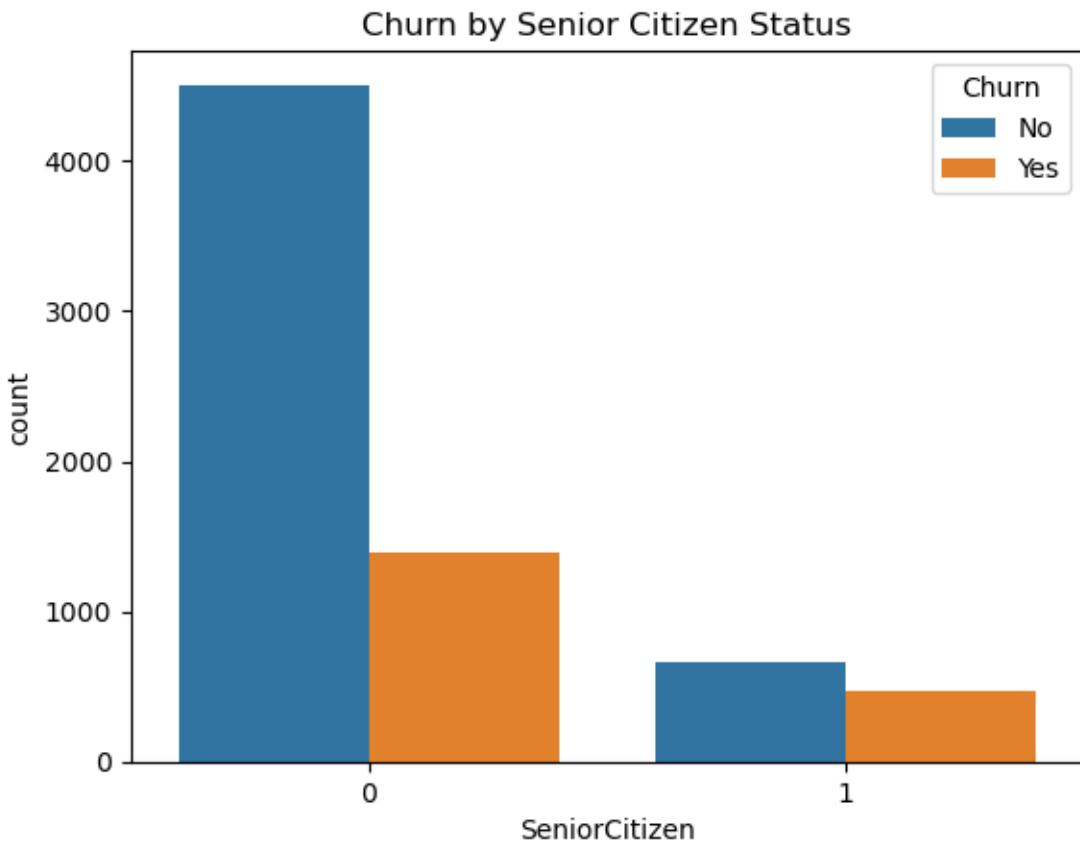
The dataset shows class imbalance, with more customers staying than leaving. This is important for modeling decisions later.

```
[19]: # Churn by Gender
sns.countplot(data=telecom_data, x='gender', hue='Churn')
plt.title("Churn by Gender")
plt.show()
```



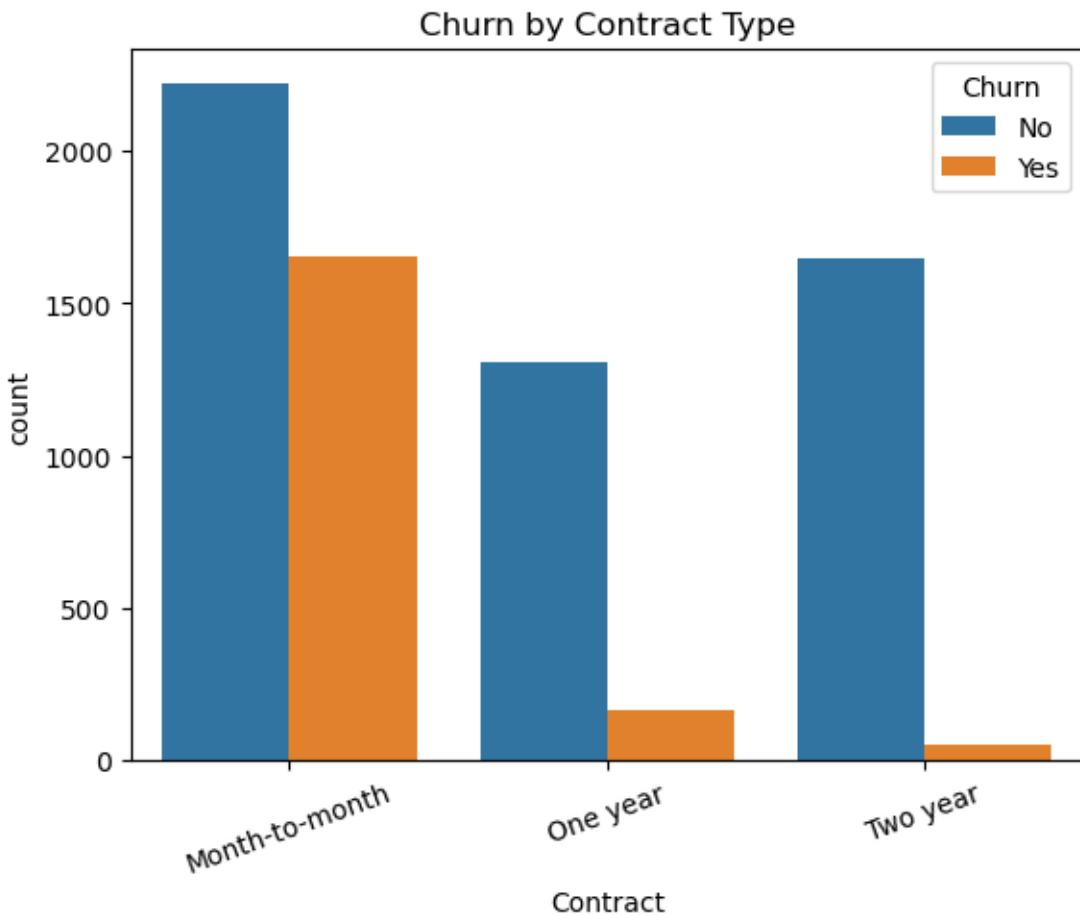
Churn rates appear similar across genders, suggesting gender alone is not a strong predictor.

```
[20]: # Churn by Senior Citizen Status
sns.countplot(data=telecom_data, x='SeniorCitizen', hue='Churn')
plt.title("Churn by Senior Citizen Status")
plt.show()
```



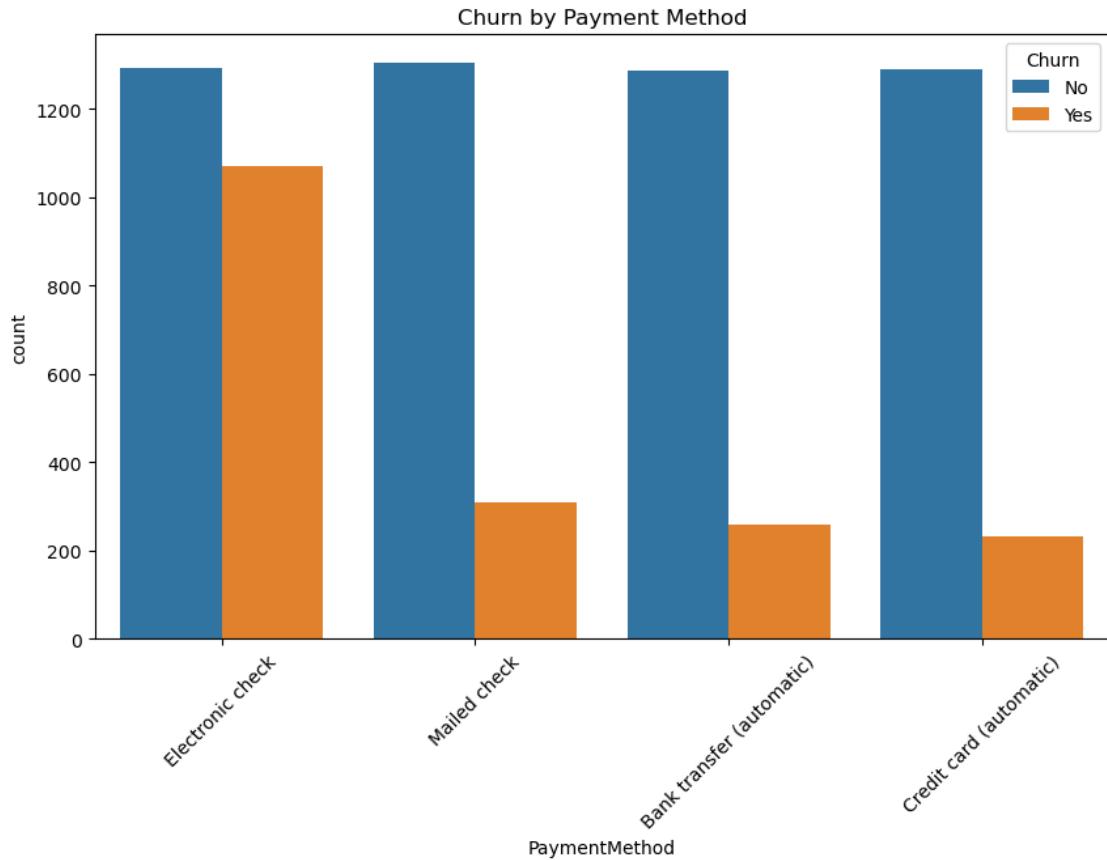
Senior citizens show a noticeably higher churn rate compared to non-senior customers.

```
[21]: # Churn by contract type
sns.countplot(data=telecom_data, x='Contract', hue='Churn')
plt.xticks(rotation=20)
plt.title("Churn by Contract Type")
plt.show()
```



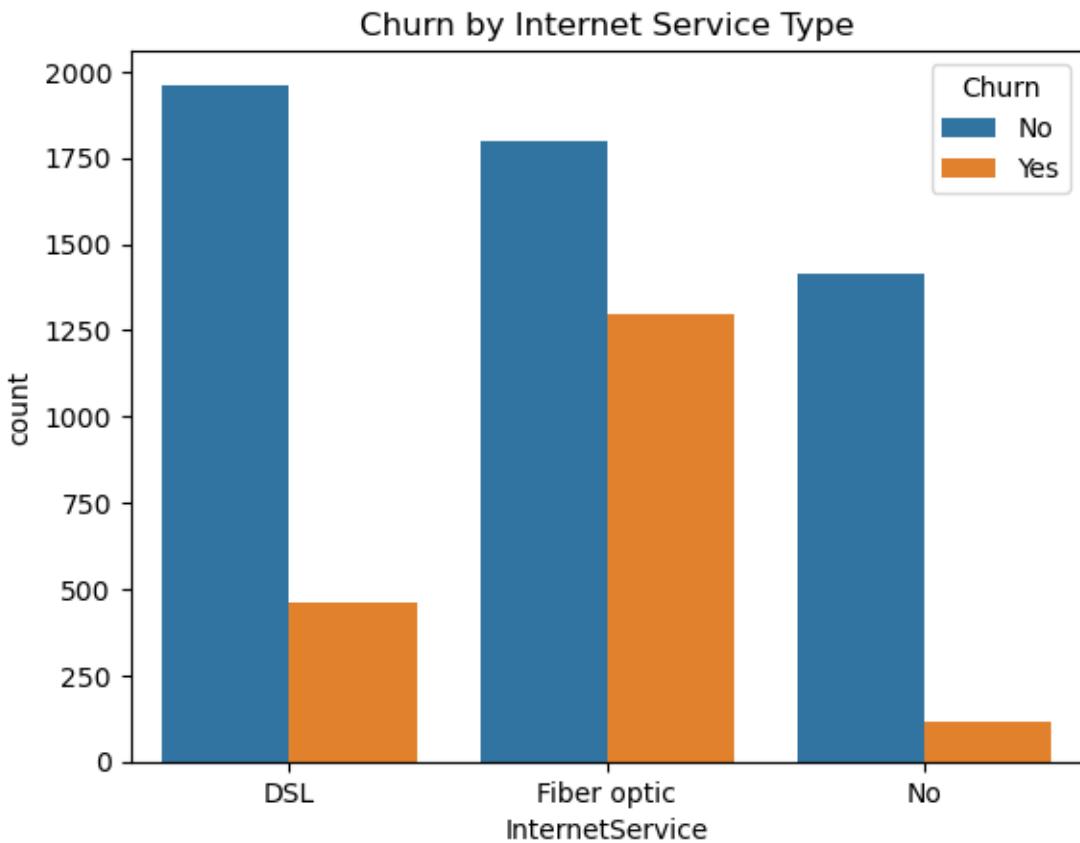
Month-to-month contracts have significantly higher churn compared to one-year and two-year contracts.

```
[22]: # Churn by Payment Method
plt.figure(figsize=(10,6))
sns.countplot(data=telecom_data, x='PaymentMethod', hue='Churn')
plt.xticks(rotation=45)
plt.title("Churn by Payment Method")
plt.show()
```

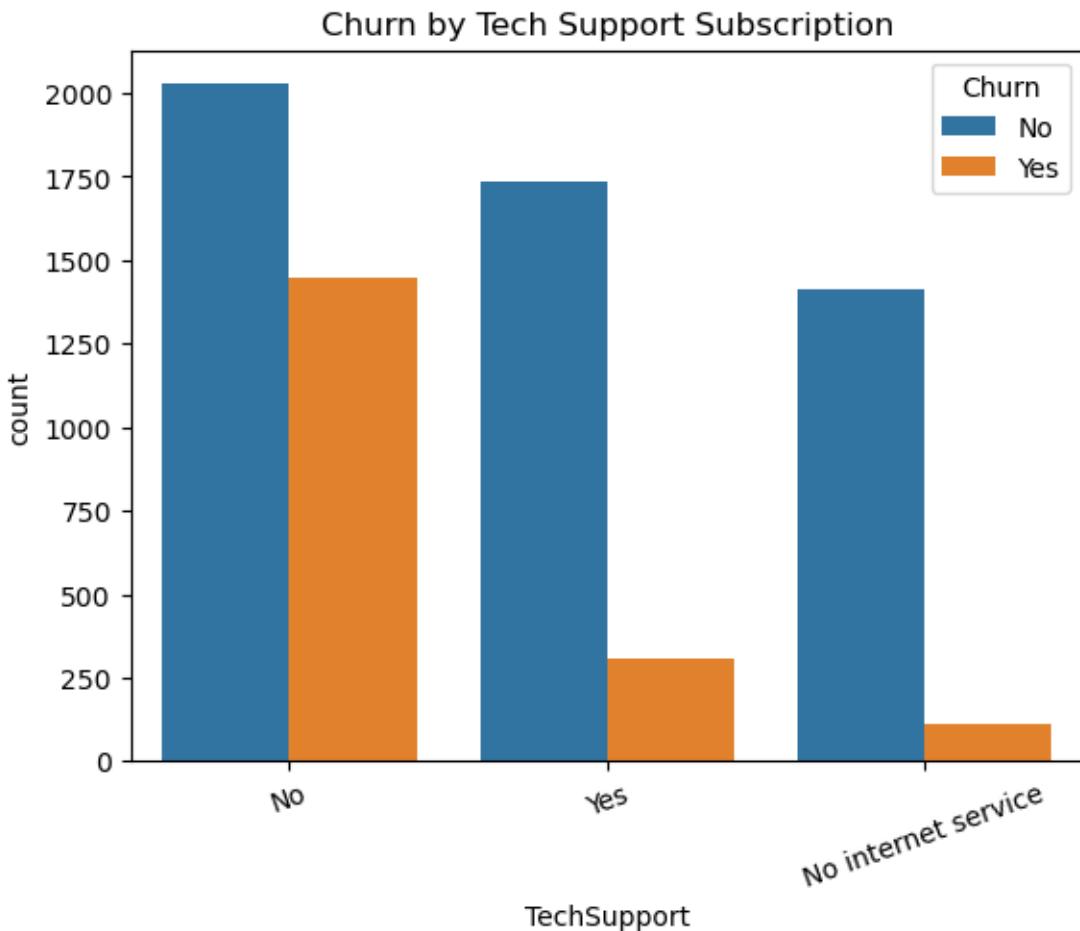


Customers using electronic checks churn more frequently than those using automatic payments.

```
[23]: sns.countplot(data=telecom_data, x='InternetService', hue='Churn')
plt.title("Churn by Internet Service Type")
plt.show()
```

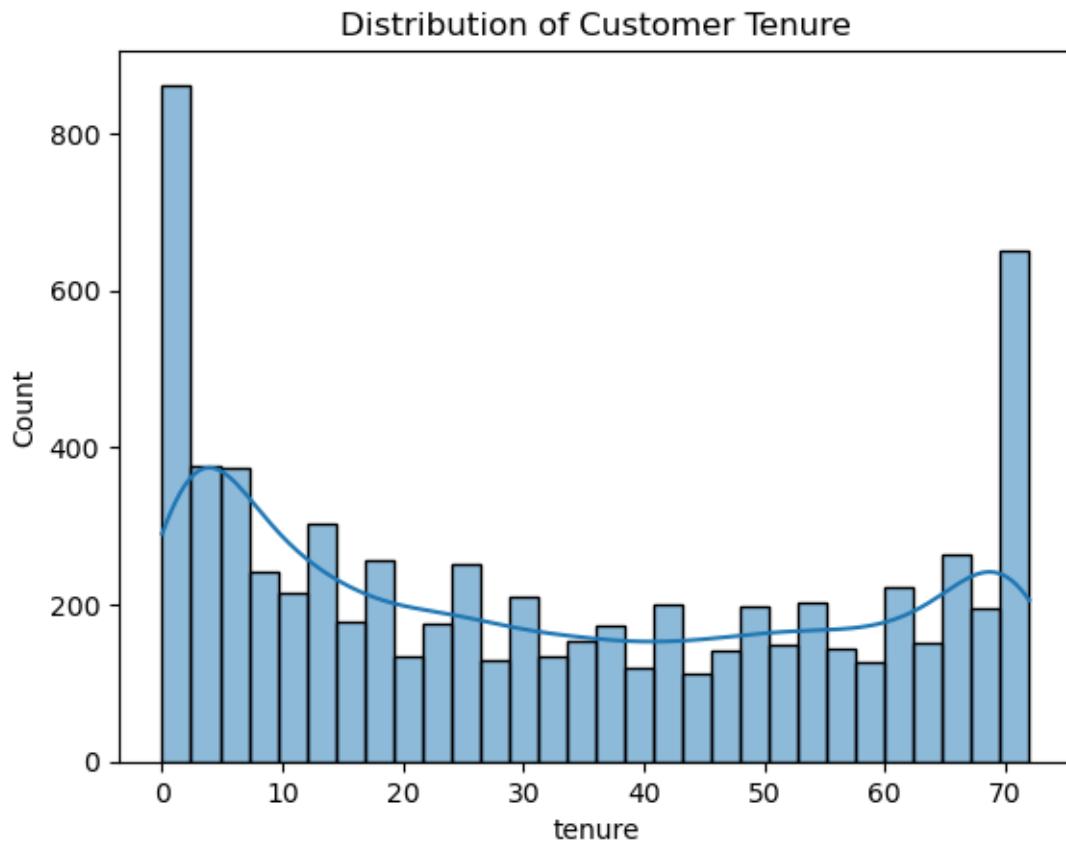


```
[24]: sns.countplot(data=telecom_data, x='TechSupport', hue='Churn')
plt.xticks(rotation=20)
plt.title("Churn by Tech Support Subscription")
plt.show()
```

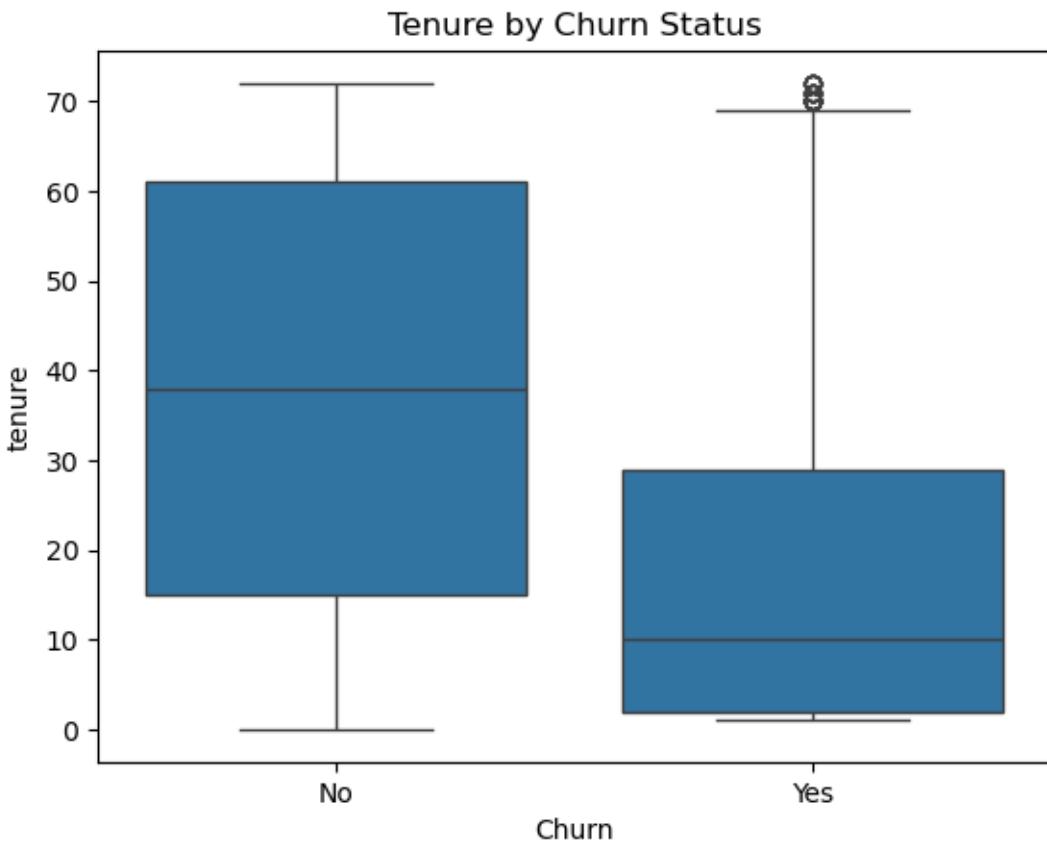


Customers without tech support churn at a much higher rate, suggesting support services improve retention.

```
[25]: sns.histplot(telecom_data['tenure'], bins=30, kde=True)
plt.title("Distribution of Customer Tenure")
plt.show()
```

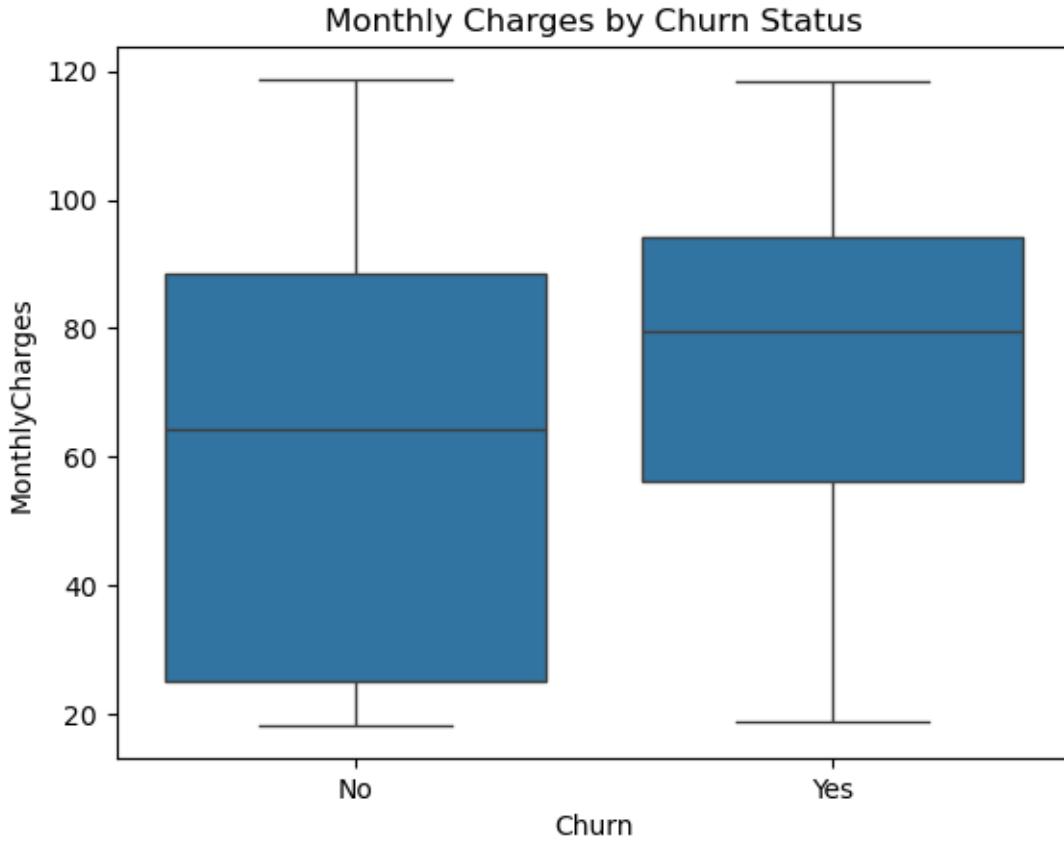


```
[26]: sns.boxplot(data=telecom_data, x='Churn', y='tenure')
plt.title("Tenure by Churn Status")
plt.show()
```



Customers who churn tend to have much shorter tenure.

```
[27]: # Monthly Charges vs Churn
sns.boxplot(data=telecom_data, x='Churn', y='MonthlyCharges')
plt.title("Monthly Charges by Churn Status")
plt.show()
```

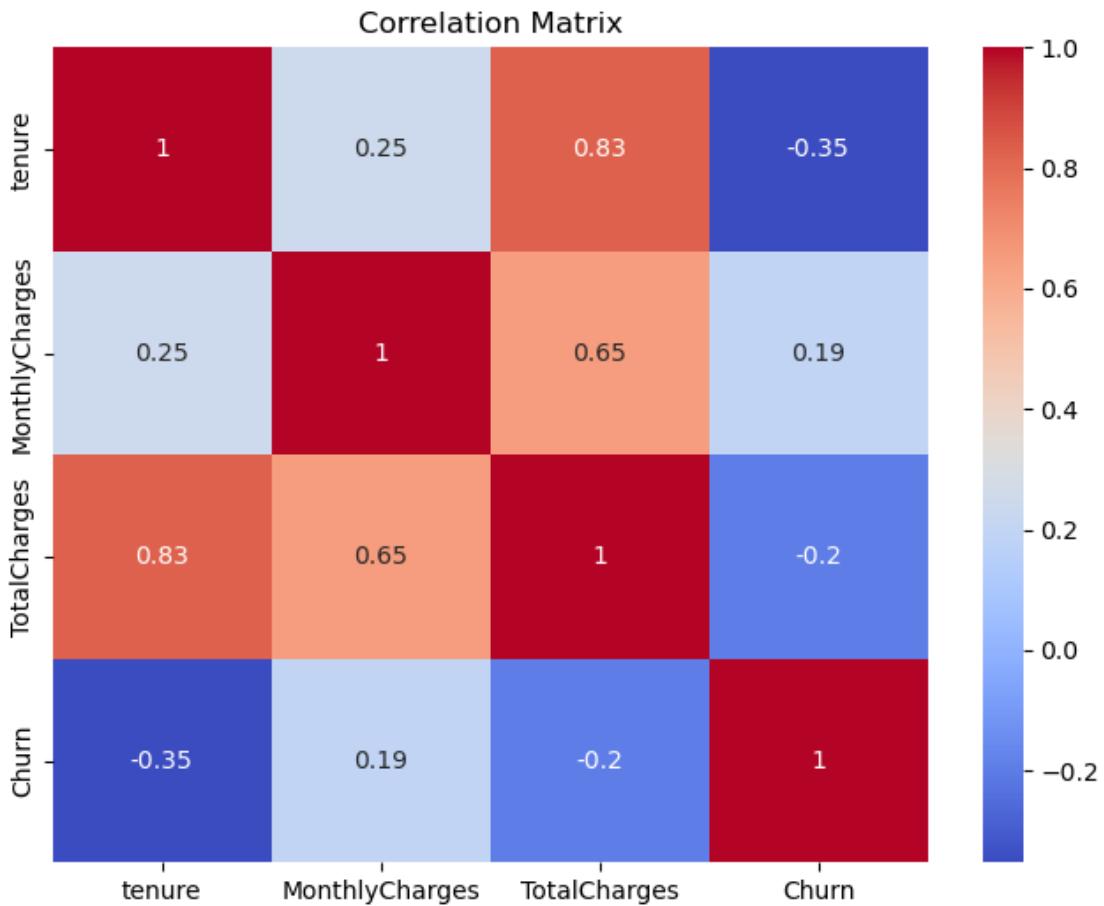


Customers with higher monthly charges are more likely to churn.

```
[28]: # Churn Correlation analysis
tc_corr = telecom_data.copy()
tc_corr['Churn'] = tc_corr['Churn'].map({'Yes': 1, 'No': 0})
```

```
[29]: plt.figure(figsize=(8,6))
sns.heatmap(tc_corr[['tenure', 'MonthlyCharges', 'TotalCharges', 'Churn']].corr(),
            annot=True, cmap='coolwarm')
plt.title("Correlation Matrix")
plt.show()
```



**Key Findings:** Tenure is strongly negatively correlated with churn. Monthly charges show a positive correlation with churn

**Key Takeaways**

- \* Customers on month-to-month contracts are the most likely to churn
- \* Short-tenure customers are significantly more likely to leave
- \* Higher monthly charges are associated with higher churn
- \* Customers lacking support services (e.g., tech support) churn more often
- \* The dataset exhibits class imbalance, which must be addressed during modeling

#### Telco Customer Churn Prediction

```
[30]: # Strip whitespace from categorical columns
for col in telecom_data.select_dtypes(include='object').columns:
    telecom_data[col] = telecom_data[col].str.strip()
```

```
[31]: # Separate Features and Target
X = telecom_data.drop('Churn', axis=1)
y = telecom_data['Churn'].map({'Yes': 1, 'No': 0})
```

```
[32]: assert y.isna().sum() == 0, "Target variable contains NaNs"
assert X.isna().sum().sum() == 0, "Feature matrix contains NaNs"
```

```
[33]: # Train, Test, Split
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2,
    random_state=42,
    stratify=y
)

[34]: # Feature type identification
categorical_features = X.select_dtypes(include='object').columns.tolist()
numerical_features = X.select_dtypes(include=['int64', 'float64']).columns.
    tolist()

[35]: # Column transformer
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numerical_features),
        ('cat', OneHotEncoder(drop='first', handle_unknown='ignore'), categorical_features)
    ]
)

[36]: log_reg_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', LogisticRegression(
        max_iter=1000,
        class_weight='balanced',
        random_state=42
    ))
])

[37]: # Train and evaluate
log_reg_pipeline.fit(X_train, y_train)

y_pred = log_reg_pipeline.predict(X_test)
y_prob = log_reg_pipeline.predict_proba(X_test)[:, 1]

[38]: from sklearn.metrics import classification_report, roc_auc_score

print(classification_report(y_test, y_pred))
print("ROC-AUC:", roc_auc_score(y_test, y_prob))
```

	precision	recall	f1-score	support
0	0.90	0.72	0.80	1035
1	0.51	0.78	0.61	374

accuracy			0.74	1409
macro avg	0.70	0.75	0.71	1409
weighted avg	0.80	0.74	0.75	1409

ROC-AUC: 0.8417499806246609

Overall Performance \* Accuracy: 74% \* ROC-AUC: 0.84 → Strong discriminative power

Even though accuracy is moderate, the ROC-AUC indicates the model is very good at ranking customers by churn risk.

```
[40]: fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = roc_auc_score(y_test, y_prob)
```

```
[ ]: # Plot the ROC Curve
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, label=f'Logistic Regression (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Customer Churn Prediction')
plt.legend(loc='lower right')
plt.tight_layout()
plt.show()
```

